

Advanced OpenMP

Advanced clauses, variable types, race conditions, task parallelism

OpenMP variable types

- When threads are created, an existing variable **x** in the master thread can have different status in the newly created threads, among which:
- **private(x)**: create private copy in each thread, uninitialized
- **firstprivate(x)**: create private copy in each thread, copy initial value
- **threadprivate(x)**: as if having a **global** copy of a variable per thread
- **shared(x)**: x of the master is shared among all threads for read/write.
- **default(shared | none)**: Variables **shared** by default, **none** disables it.
 - Using **none** is a good practice for preventing careless bugs
- **reduction(op:x)**: Creates a safe copy of a **shared x** for each thread, merges them transparently at the end using **op**. Avoids race conditions.

Example: OpenMP variable types

- `openmp-example-variables.cpp`
- `openmp-example-variables-2.cpp`

nowait clause

- Removes the **implicit barrier** at the end of a work sharing construct
 - Enables threads to continue working in the rest of the parallel region
- Can be added at the end of **omp for**, **omp sections**, **omp single**
- Should be added diligently; make sure that there is no dependency on the work sharing construct in the rest of the code

```
int N = ...;
#pragma omp parallel default(none) num_threads(P)
shared(N)
{
    #pragma omp for nowait
    for (int i = 0; i < N; i++) {
        f(i);
    } // end of for, no barrier due to nowait

    // Threads can start the second loop immediately after
    // finishing their chunk for the first loop
    #pragma omp for
    for (int i = 0; i < N; i++) {
        g(i);
    } // end of for, implicit barrier
}
```

if clause

```
#pragma omp parallel num_threads(P) if(cond)
{
    // Parallel code to be executed by each thread
}
```

- **if** clause can be added when creating a parallel region or a task
- Threads/tasks are created only if the given **cond** is true/nonzero
 - Otherwise, only the main thread executes the block of code
- Useful for preventing thread/task creation overhead for small problems
 - Example:
 - **#pragma omp parallel if (N > 128)**
 - **#pragma omp task if (N > 256)**

schedule clause for omp for

- Determines how iterations are distributed among threads
- **schedule(policy, chunksize)**
- **static** policy: Assign chunksize contiguous iterations to each thread in a circular order (0,1,2,3,0,1,2,3,...)
- **dynamic** policy: Assign chunksize contiguous iterations to the first available thread (0,3,2,0,1,2,1,3,2,0,...)
- **guided** policy: Start assigning big chunks dynamically, gradually reduce it to chunksize towards the end.
- **runtime** policy: Defer the decision to user in runtime (e.g. using the OMP_SCHEDULE variable)

```
int N = ...;
#pragma omp parallel default(none) num_threads(P)
shared(N)
{
    #pragma omp for schedule(dynamic, 32)
    for (int i = 0; i < N; i++) {
        f(i);
    } // end of for, implicit barrier
    #pragma omp for schedule(static)
    for (int i = 0; i < N; i++) {
        g(i);
    } // end of for, implicit barrier
}
```

static loop scheduling

- Specified with the clause **schedule**(static, chunksize)
- **static** schedule
 - Gives chunksize contiguous iterations to threads in a round-robin manner (th0, th1, ..., thP-1, th0, th1, ...)
 - If no chunksize is given, it is N/P by default (single big chunk per thread)
 - If no schedule is specified, it is static with chunksize=N/P by default
 - Good strategy if the workload per iteration is uniform

dynamic loop scheduling

- Specified with the clause **schedule(dynamic, chunksize)**
- **dynamic** schedule
 - Gives chunksize contiguous iterations to the first available thread (th0, th1, th0, th3, th1, ...)
 - If no chunksize specified, it is 1 by default (single iteration per chunk)
 - Good strategy when the workload among iterations vary significantly
 - **dynamic** makes threads idle less due to better workload balance
 - Potentially bad for data locality per thread
 - Each thread might touch many non-contiguous chunks in an array.
 - If chunksize is big enough (giving ≥ 64 byte chunks, e.g. 16 ints), this is less of a problem.

Controlling loop scheduling in runtime

- If no schedule policy is specified in a loop with a **schedule** clause,
 - Setting OMP_SCHEDULE variable to "schedule,chunksize" modifies the loop scheduling in runtime
 - Ex: OMP_SCHEDULE="guided,4" ./prog
 - omp_set_schedule(policy, chunksize) can also modify the scheduling
 - Ex: omp_set_schedule(omp_sched_dynamic);

Example: OpenMP loop scheduling

- `openmp-example-schedule.cpp`

Merging **parallel** and **for/sections** directives

- If the parallel region has an **omp for** or **omp sections** inside, and **nothing else**, you can merge two directives.
- You cannot use curly braces (**{ ... }**) after **#pragma omp parallel for**.
- You cannot put anything else other than **section** blocks after **#pragma omp parallel sections**

```
#pragma omp parallel num_threads(P)
{
    #pragma omp for
    for (int i = 0; i < N; i++) {
        f(i)
    }
}
```

```
#pragma omp parallel for num_threads(P)
for (int i = 0; i < N; i++) {
    f(i)
}
```

```
#pragma omp parallel sections num_threads(P)
{
    #pragma omp section
    {
        g();
    }
    #pragma omp section
    {
        h();
    }
}
```

collapse clause for omp for

- **collapse** clause flattens nested loops into a single domain to provide more parallelism (e.g., when iteration domain is small).
- Iteration domain must be a fixed size for all nested loops.
- Can collapse multiple loops (≥ 2)

```
#pragma omp parallel num_threads(P)
{
    // This loop is
    #pragma omp for collapse(2)
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            f(i, j);
        }
    }
}
```

```
// Same as this
#pragma omp for
for (int i = 0; i < N*N; i++) {
    f(i/N, i%N);
}
}
```

ordered clause for omp for

- Specifies a block in a parallel loop that respects the sequential order of iterations in execution.
- Useful for gathering results in order at the end, after performing expensive computations in parallel.
- Works for all loop schedules.

```
// We want to have result[i] = f(i) in the end  
std::vector<int> result;
```

```
#pragma omp parallel num_threads(P)  
{  
#pragma omp for ordered  
    for (int i = 0; i < N; i++) {  
        int res = f(i); // Done in parallel, time consuming  
#pragma omp ordered  
        result.push_back(res); // Done in order sequentially  
    }  
}
```

Race conditions

- What is the final value of n?
min/max?
- There are three hidden instructions:
 - **load** n
 - **add** n, 1
 - **store** n
- Threads simultaneously read the old value before updating n
 - Some increments are lost
 - => **race condition**
- Solutions?
 - **atomic**
 - **critical**
 - **reduction**

```
// We want to have result[i] = f(i) in the end  
int n = 0;
```

```
#pragma omp parallel for num_threads(P)  
  for (int i = 0; i < 1000000; i++) {  
    n++;  
  }
```

Race conditions

- What is the final value of n? min/max?
- There are three hidden instructions:
 - **load** n
 - **add** n, 1
 - **store** n
- Threads simultaneously read the old value before updating n
 - Some increments are lost
 - => **race condition**
- Solutions?
 - **atomic** directive
 - Uses hardware atomic instructions for the following instruction, has some overhead
 - Works for basic arithmetic / logical operations (+, -, *, /, min, max, &, |, ...)

```
// We want to have result[i] = f(i) in the end  
int n = 0;
```

```
#pragma omp parallel for num_threads(P)  
    for (int i = 0; i < 1000000; i++) {  
#pragma omp atomic  
        n++;  
    }
```

Race conditions

- What is the final value of n? min/max?
- There are three hidden instructions:
 - **load** n
 - **add** n, 1
 - **store** n
- Threads simultaneously read the old value before updating n
 - Some increments are lost
 - => **race condition**
- Solutions?
 - **critical** directive
 - Block of code in a **critical** region is executed by a **single thread at a time**
 - Works for all code types but higher overhead than **atomic**

```
// We want to have result[i] = f(i) in the end  
int n = 0;
```

```
#pragma omp parallel for num_threads(P)  
    for (int i = 0; i < 1000000; i++) {  
#pragma omp critical  
    {  
        n++;  
    }  
}
```


Race conditions

- What is the final value of n? min/max?
- There are three hidden instructions:
 - **load** n
 - **add** n, 1
 - **store** n
- Threads simultaneously read the old value before updating n
 - Some increments are lost
 - => **race condition**
- Solutions?
 - **reduction(op:x)** clause
 - Creates an invisible copy of the shared variable **x** in each thread, performs the reduction operation **op** on the private copy without atomic nor critical.
 - At the end of the loop, merges these private copies with a single atomic operation **op** on the shared **x**.
 - **op** can be +, -, *, /, min, max, &, |, &&, | |, ^,

```
// We want to have result[i] = f(i) in the end  
int n = 0;
```


```
#pragma omp parallel for num_threads(P) reduction(+:n)  
    for (int i = 0; i < 1000000; i++) {  
        n++;  
    }  
}
```

OpenMP task parallelism

Task creation, management, and dependencies

Task-based parallelism

- A task is a block of computation (block of code, function, even loop iterations)
- A task might need another task's output, which is called a **dependency**.
- Task-based parallelism expresses the computation as a bunch of tasks, and their interdependencies.
- A **runtime system** then orchestrates their execution among multiple threads.
- It is particularly useful for
 - irregular parallel computations
 - recursive functions
- Advanced task-based runtime systems enable hybrid GPU+CPU+FPGA parallelism
 - For each task, specify data dependencies
 - For each device/task, specify code/kernel
 - Runtime system orchestrates the distribution of tasks to compute units



Creating tasks in OpenMP

- When a thread encounters a **#pragma omp task** directive, it creates a task involving the following code block, and puts it into a task pool instead of executing it.
- When one of the threads in the parallel region is idle (i.e., stuck at a barrier)
 - it executes one of the **ready** tasks (whose all dependencies are satisfied) in the pool.
- No particular order for task execution, aside from dependencies.
- The program on the right, is it correct?

```
#pragma omp parallel num_threads(P)
{
    #pragma omp task
    f();
    #pragma omp task
    g();
}
```

Creating tasks in OpenMP

- Each task must be created by a **single thread**
 - In many cases, need to use a work sharing construct (**omp single/master** or **omp for**)
 - Generally not needed in recursive functions
- **#pragma omp taskwait** blocks the thread until all tasks created by the thread are done
 - It **does not** wait for all tasks globally
- **#pragma omp barrier** implies a **taskwait**
 - Therefore, at the end of **omp single/for**, all tasks are guaranteed to be executed.
- The program on the right, is it correct?

```
#pragma omp parallel num_threads(P)
{
    #pragma omp single
    {
        #pragma omp task
        f();
        #pragma omp task
        g();
        #pragma omp taskwait
        // After taskwait f() and g() are guaranteed to
        finish.
    }
    // Here, even without taskwait f() and g() are executed.
}

#pragma omp parallel for num_threads(P)
for (int i = 0; i < N; i++) {
    #pragma omp task
    h(i);
}
```

Task dependencies

- Tasks can have dependencies on variables
- Dependency can be
 - **input** dependency:
 - **depend(in:var)** clause after **task**
 - **output** dependency:
 - **depend(out:var)** clause after **task**
- A task with **depend(in:x)** can start only after **all previously submitted tasks** with **depend(out:x)** finish.
- Sequential order of task creation and their input/output dependencies determine the order of execution. For example, if task t2 is submitted after t1, then
 - If t2 has **depend(in:x)** and t1 has **depend(in:x)**, both can be executed in any order
 - If t2 has **depend(in:x)** and t1 has **depend(out:x)**, t2 must wait for t1 to terminate
 - If t2 has **depend(out:x)** and t1 has **depend(in:x)**, then t2 must wait for t1 to terminate
 - If t2 has **depend(out:x)** and t1 has **depend(out:x)**, then t2 must wait for t1 to terminate
- It is the programmer's responsibility to correctly choose task creation order and dependencies.

```
// We want to compute x = f() + g();
int x;
int y;

#pragma omp parallel num_threads(P)
{
    #pragma omp single
    {
        #pragma omp task depend(out:x)
        x = f();
        #pragma omp task depend(out:y)
        y = g();
        #pragma omp task depend(in:x) depend(in:y)
        x = x + y;
    }
    #pragma omp taskwait
    printf("f() + g() = %d\n", x);
}
```

Task dependencies

- Dependencies can be on array elements
- If tasks work on blocks of data (continuous subarray, submatrix) no need to declare dependency for each data element
 - declare a dependency on **the same** element of the block (e.g., first element of the subarray, upper-left element of the submatrix) **across all tasks**

```
// We want to compute  $x = f() + g()$ ;  
int x[3];  
  
#pragma omp parallel num_threads(P)  
{  
    #pragma omp single  
    {  
        #pragma omp task depend(out:x[0])  
        x[0] = f();  
        #pragma omp task depend(out:x[1])  
        x[1] = g();  
        #pragma omp task depend(in:x[0]) depend(in:x[1])  
        depend(out:x[2])  
        x[2] = x[0] + x[1];  
        #pragma omp taskwait  
        printf("f() + g() = %d\n", x[2]);  
    }  
}
```

Example: Computing nine

- `openmp-example-tasks.cpp`
- `openmp-example-tasks-deps.cpp`

Task variable visibility

- Shared variables in the **enclosing scope** stay shared with the task created in that scope
- Private variables in the **enclosing scope** become **firstprivate** in the task created in that scope
- A variable declared within the task is **private** to that task, and is destroyed once the task is done

```
int x = 0;

#pragma omp parallel num_threads(P) // x shared by default
{
    #pragma omp single
    {
        int y = 0;
        #pragma omp task
        {
            // task has shared access to x
            x = 1;
            // task has its firstprivate copy of y = 0
            y = 2; // y in the upper scope is not modified
        } // firstprivate y of the task gets destroyed
    }
    #pragma omp taskwait
    printf("x = %d\n", x); // x is modified, prints 1
    printf("y = %d\n", y); // y is unmodified, prints 0
}
```